

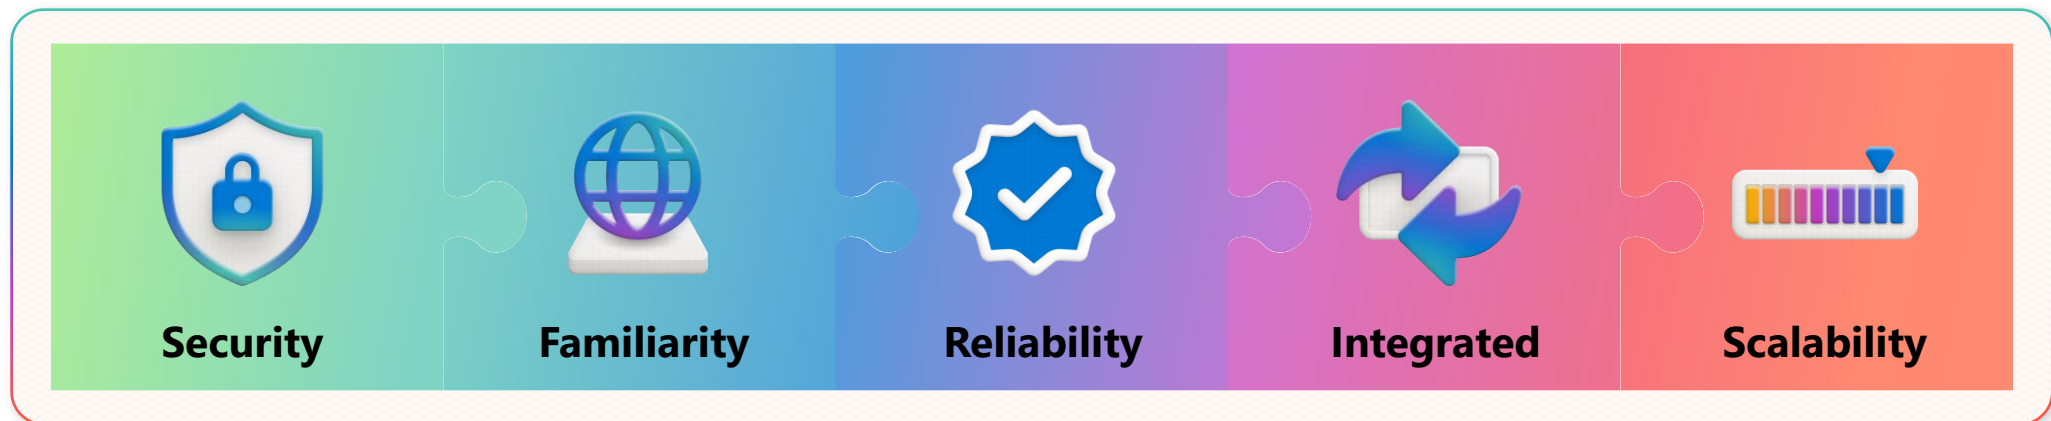


# Leveraging AI capabilities in your modern data workload



# What problems are we trying to solve with AI?

- ✓ *Smarter searching* on your existing data
- ✓ Bring in other documents/text for *centralized vector searching*
- ✓ Provide *building blocks* for Intelligent assistants, RAG, AI Agents
- ✓ Take advantage of AI in a *secure and scalable* fashion
- ✓ *Overcome complexity* by using the familiar SQL language





# Vector Search in Azure Cosmos DB



# Flexible schema

SELECT \* FROM c

Type a query predicate (e.g., WHERE c.id='1'), or choose one from the drop down list, or leave empty to query a

Home

vector-nosql-db

vector-diskann

Items

Scale & Settings

Stored Procedures

User Defined Functions

Triggers

vector-nosql-cont

Items

Scale & Settings

Stored Procedures

User Defined Functions

<input type="checkbox"/>	id	...	/partKey	...	
<input checked="" type="checkbox"/>	s1		2020		
<input type="checkbox"/>	s2		2021		
<input type="checkbox"/>	s3		2021		
<input type="checkbox"/>	s4		2021		
<input type="checkbox"/>	s5		2021		
<input type="checkbox"/>	s6		2021		
<input type="checkbox"/>	s7		2021		
<input type="checkbox"/>	s8		1993		
<input type="checkbox"/>	s9		2021		
<input type="checkbox"/>	s10		2021		
<input type="checkbox"/>	s11		2021		
<input type="checkbox"/>	s12		2021		

1

{

"id": "s1",

"type": "Movie",

"title": "Dick Johnson Is Dead",

"director": "Kirsten Johnson",

"cast": "",

"country": "United States",

"date\_added": "September 25, 2021",

"release\_year": "2020",

"rating": "PG-13",

"duration": "90 min",

"description": "As her father nears the end of his life

"listed\_in": "Documentaries",

"docVector": [

-0.010484142228960991,

0.006622579880058765,

-0.00941577646881342,

-0.004073948599398136,

0.009344981051981449,

0.022101009264588356,

-0.013110005296766758,

-0.013644187711179256,

-0.006912196986377239,

-0.006139884702861309,

0.003584817284718156,

0.021238593384623528.

}

26



# Indexing

- Create a DiskANN **index on the document embedding path**
- **Add the vector path to the “excludedPaths”** section of the indexing policy (to avoid indexing as a regular array path)
- Skip vector indexing if type and dimensions differ from container vector policy
- **DiskANN requires at least 1000 documents**, otherwise will search all documents
- **Best practice:** exclude everything and selective add properties to be indexed

```
1  {
2      "indexingMode": "consistent",
3      "automatic": true,
4      "includedPaths": [
5          {
6              "path": "/*"
7          }
8      ],
9      "excludedPaths": [
10         {
11             "path": "/docVector/*"
12         },
13         {
14             "path": "/\"_etag\"/?"
15         }
16     ],
17     "vectorIndexes": [
18         {
19             "path": "/docVector",
20             "type": "diskANN"
21         }
22     ]
23 }
```

# Indexing

- More than one vector policy/index is supported

The screenshot displays a 'New Container' configuration window with two vector embedding policies. The first policy, 'Vector embedding 1', is configured with the following settings: Path (/vector1), Data type (float32), Distance function (cosine), Dimensions (1536), Index type (diskANN), Quantization byte size (empty), and Indexing search list size (100). The second policy, 'Vector embedding 2', is partially visible and configured with: Path (/vector2), Data type (float32), Distance function (cosine), Dimensions (1536), and Index type (none). A dropdown menu for the 'Index type' of 'Vector embedding 2' is open, showing options: none, flat, diskANN, and quantizedFlat. The 'Add' button is partially visible at the bottom left.

**New Container**

**Vector embedding 1**

Path: /vector1

Data type: float32

Distance function: cosine

Dimensions: 1536

Index type: diskANN

Quantization byte size:

Indexing search list size: 100

**Vector embedding 2**

Path: /vector2

Data type: float32

Distance function: cosine

Dimensions: 1536

Index type: none

flat

diskANN

quantizedFlat

Add

# Ingestion

- <https://devblogs.microsoft.com/cosmosdb/azure-cosmos-db-vector-search-with-diskann-part-1-full-space-search/>
- 20 RUs for a small document with a 768-dimensional vector
- 30 RUs for a small document with a 1,536-dimensional vector
- 50 RUs for a small document with a 3,072-dimensional vector
- Document size and regular indexing policy will add to the base cost

# Latency

Scenario	P50 Latency (ms) k=10/50	P95 Latency (ms) k=10/50	Avg Latency (ms) k=10/50	Recall@k k=10/50
100k vectors	25/95	29/108	25/95	92.47/98.03
1M vectors	34/130	51/164	39/133	89.2/95.41
35M vectors	108/544	166/879	112/569	90.73/95.67



# Query cost (RU)

Scenario	P50 RU Cost k=10/50	P95 RU Cost k=10/50	Avg RU Cost k=10/50
100k vectors	36/159	39/169	36/159
1M vectors	38/170	42/179	45/182
35M vectors	282/1,249	300/1,298	282/1,245

# Cost estimation by provision mode/workload

**Monthly** cost (in USD) for Azure Cosmos DB for NoSQL in the East US 2 region

Scenario	Manual (100% 10 QPS)	Manual (100% 100 QPS)	Autoscale (50% 1 QPS, 50% 10 QPS)	Autoscale (100% 10QPS)	Autoscale (50% 10 QPS, 50% 100 QPS)	Autoscale (100% 100 QPS)
100k vectors	\$22.46	\$224.64	\$18.53	\$33.70	\$185.33	\$336.96
1M vectors	\$24.19	\$241.92	\$19.96	\$36.29	\$199.58	\$362.88
35M vectors	\$172.80	\$1,728.00	\$142.56	\$259.20	\$1,425.60	\$2,592.00



# Vector Search in Azure PostgreSQL





# Vector Search in Azure SQL



# Finding the Best Focaccia Bread in town

# How AI and SQL can help you?

# Focaccia is "similar" to

- Pizza
- Flatbread
- Bread

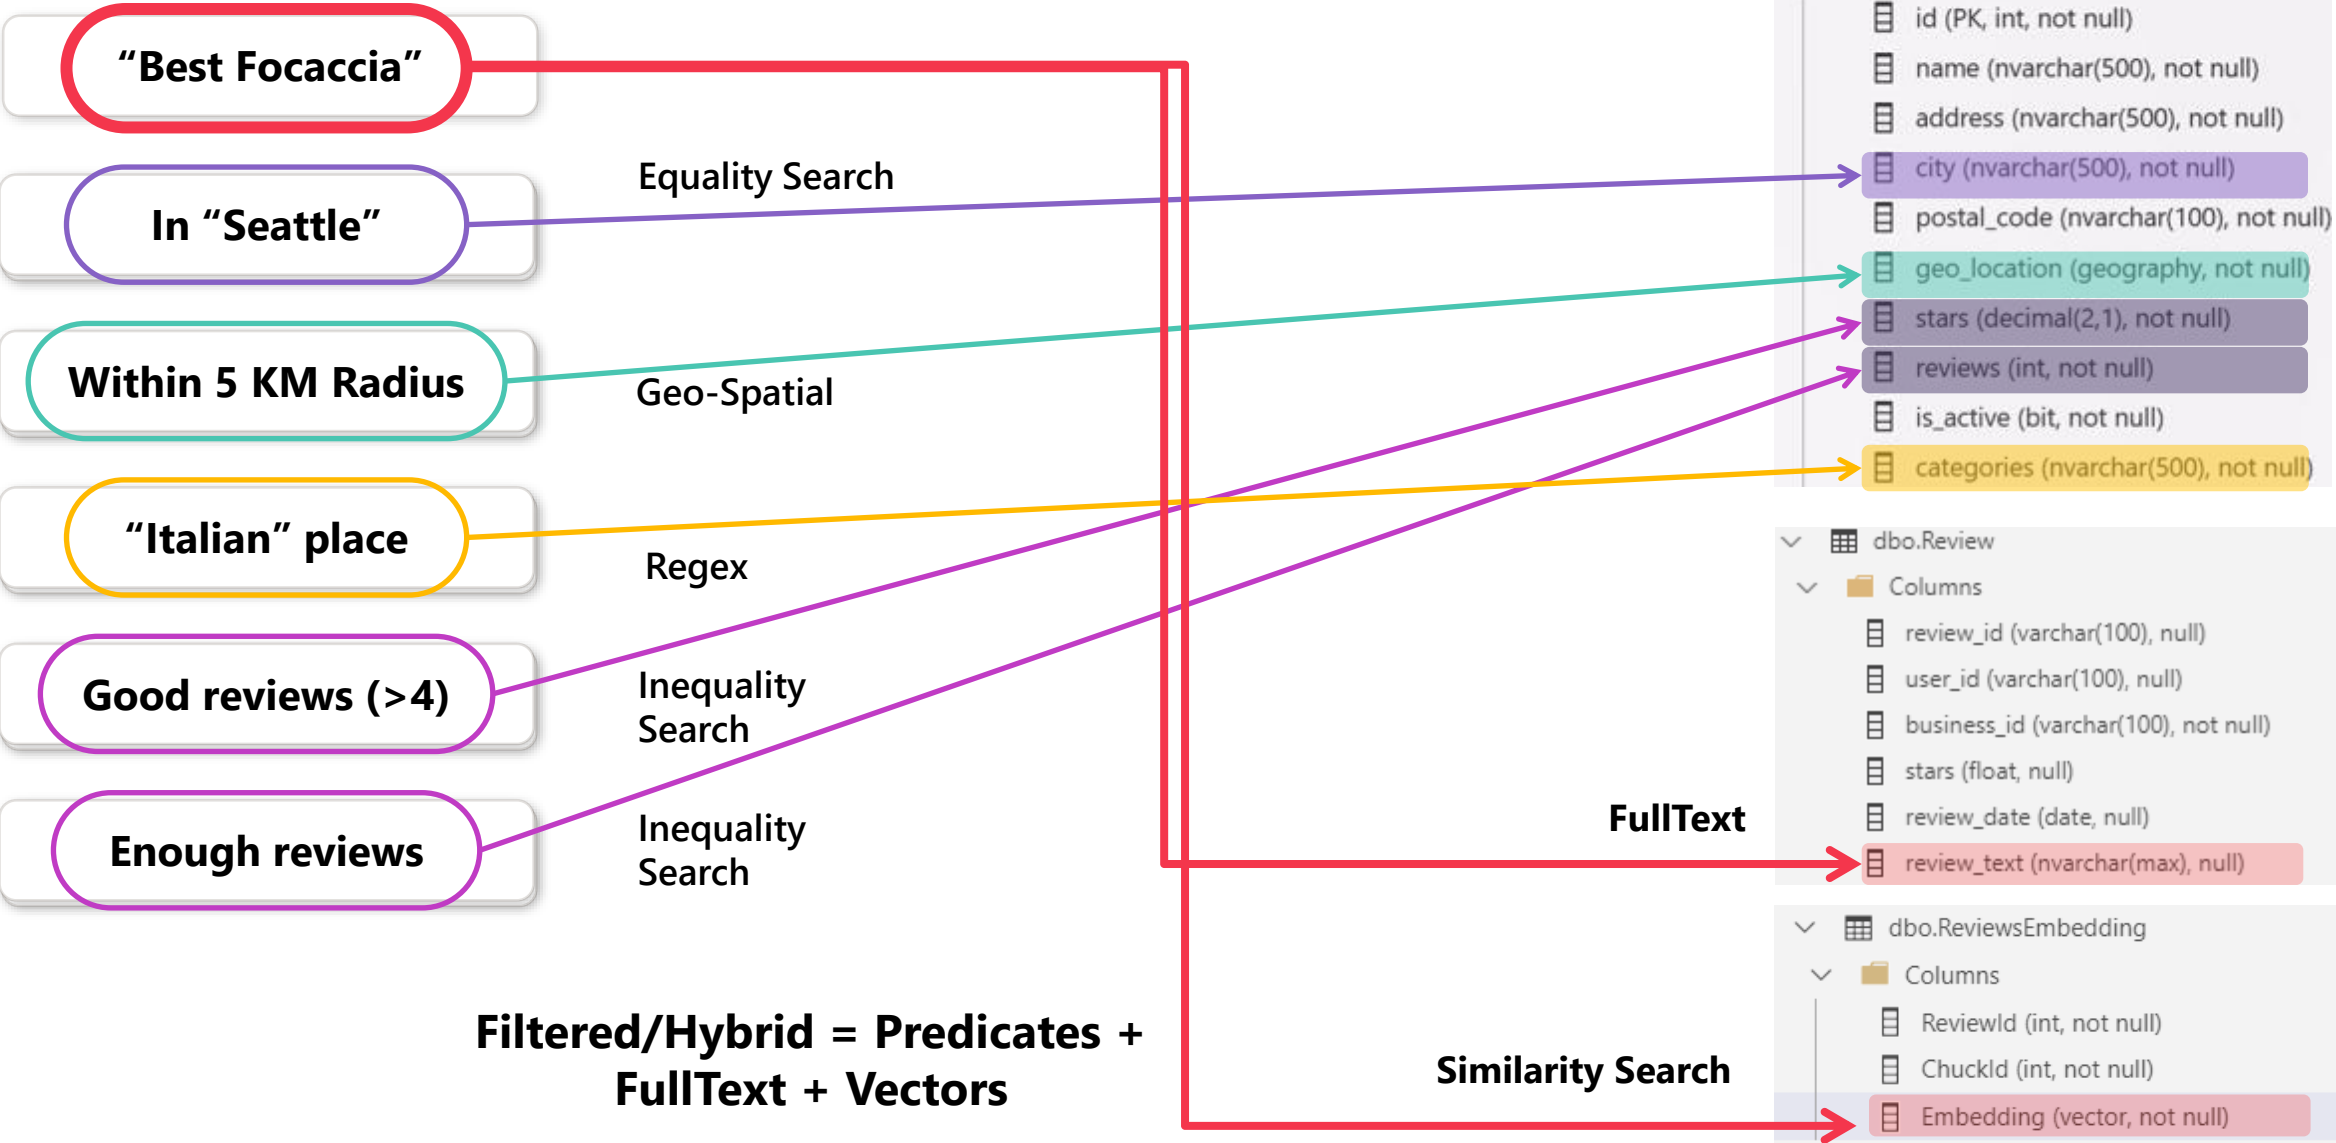
## What about misspellings?

## Focaccia's embedding:



[ -8.2512591e-003, -2.8406959e-002, 4.7846469e-003, -3.2282826e-002, -1.9243909e-003, -4.1466937e-003, -1.5021985e-002, 2.9538423e-002, 9.7979931e-003, -4.3428943e-002, -1.4350179e-004, -2.1449661e-002, -2.1991320e-002, 8.2512591e-003, -7.0716478e-003, -1.5111886e-004, 1.1447041e-002, -2.0679303e-002, -8.9900201e-004, -4.3061818e-003, 5.5802822e-002, -1.5864564e-002, 1.0688720e-002, 3.4016129e-002, -2.2424646e-002, 7.3364582e-003, -2.9105097e-002, 1.3240532e-002, 2.7564380e-002, 9.5091090e-003, 1.2482210e-002, -5.0410315e-002, 1.3473745e-003, -5.7295393e-002, 1.5783316e-003, -2.4434799e-002, -1.4624769e-003, -1.1332692e-002, -1.1868332e-002, 1.7284913e-002, -2.7010685e-002, 1.0893347e-002, 5.9124991e-002, -1.5455311e-002, 5.8450930e-002, -1.0243357e-002, 2.3652405e-002, 2.1690398e-002, -2.0920040e-002, 3.1055065e-002, 1.4432180e-002, -8.1970925e-003, -1.2229437e-002, 2.6631525e-003, -1.5202538e-002, -1.4648843e-002, 1.5936786e-002, 5.8258340e-002, -2.0270050e-002, 6.4902678e-002, -2.7179200e-002, -1.5539570e-002, 8.6857885e-002, -2.9345833e-002, -1.1908956e-003, -4.5770109e-003, -1.3627748e-002, 8.1549641e-003, -7.5290478e-003, 3.1873573e-002, -5.2119549e-003, 1.7441392e-002, -1.8633040e-002, 1.2783132e-002, 4.6967778e-002, -6.6323029e-003, 6.3627748e-002, 8.1549641e-003, -7.5290478e-003, 3.1873573e-002, -5.2119549e-003, 1.7441392e-002, -1.8633040e-002, 1.2783132e-002, 4.6967778e-002, -6.6323029e-003, 6.3627748e-002 ]

# Finding the Best Focaccia Bread in town





```

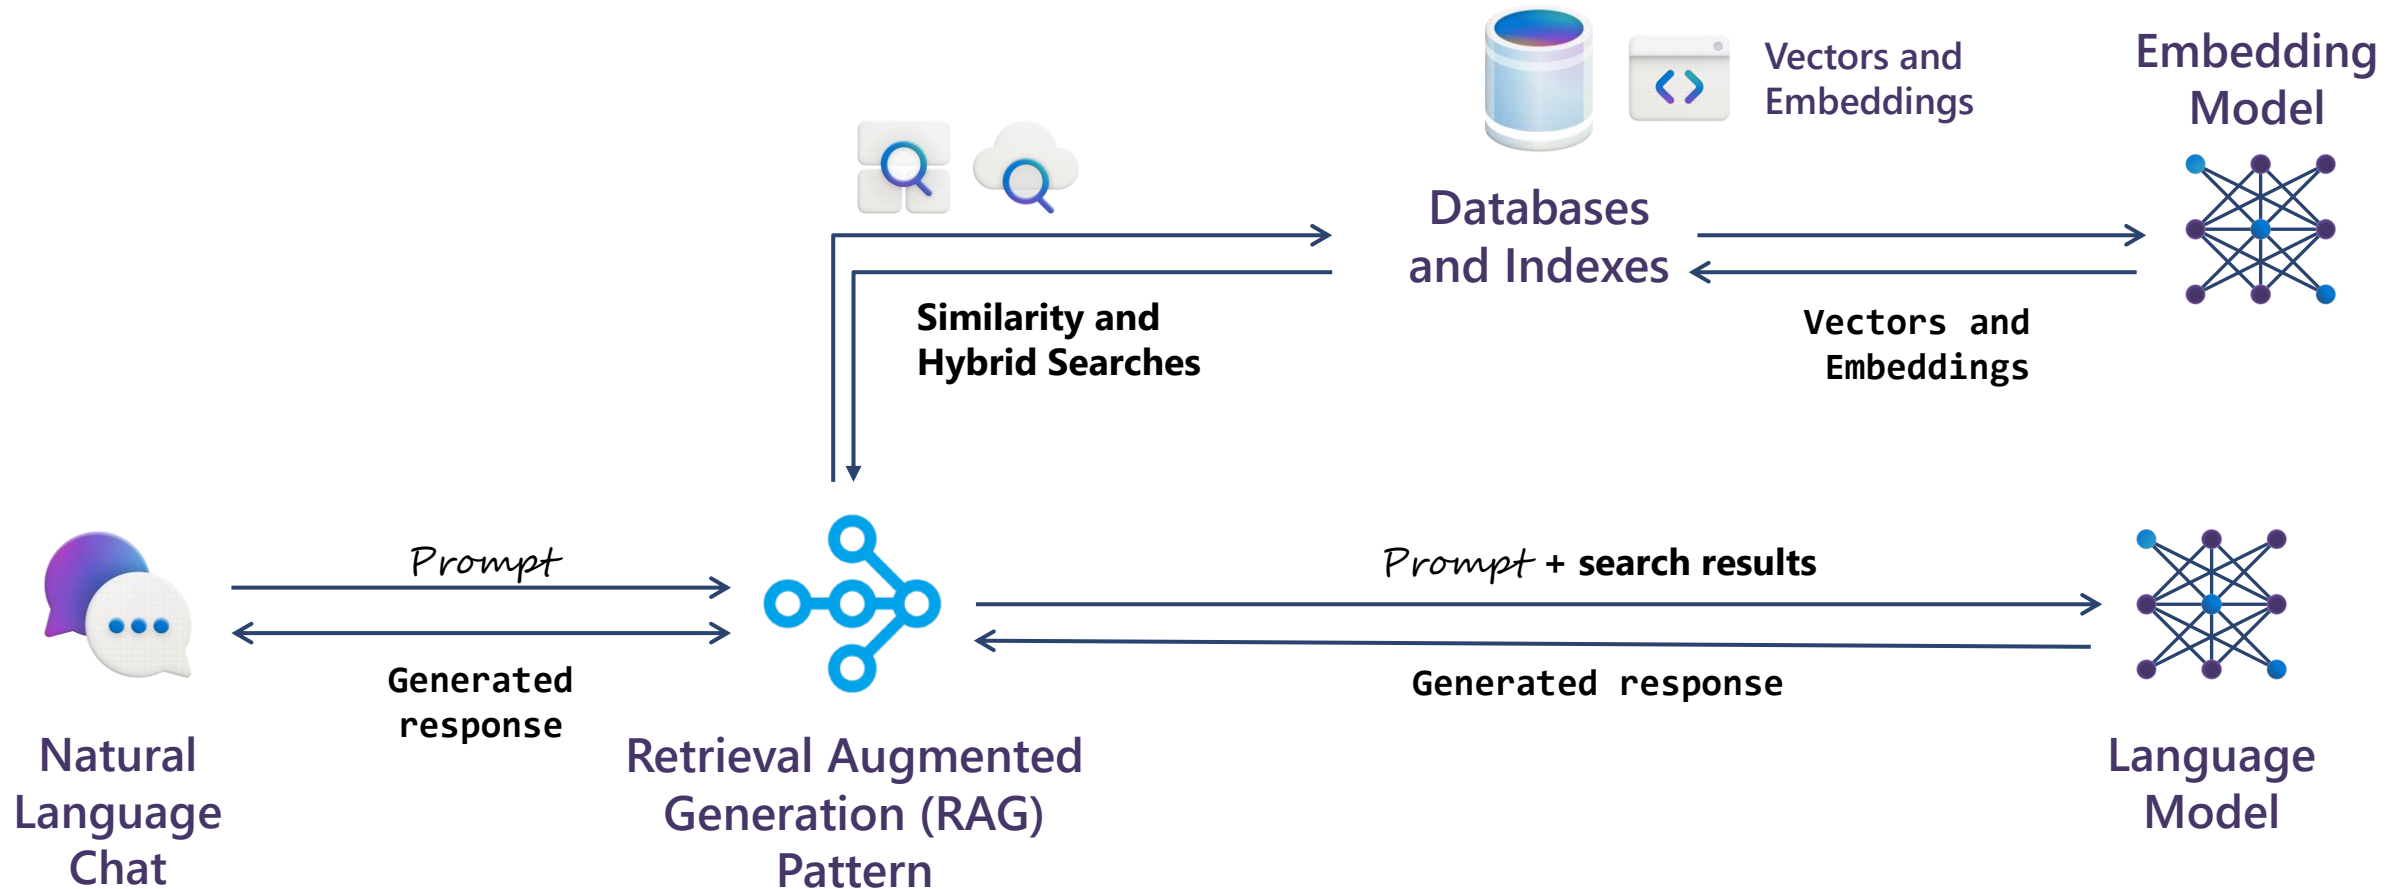
-- Complete query
select top(30)
    b.id as business_id,
    b.name as business_name,
    r.id as review_id,
    r.stars,
    r.review,
    1-vector_distance('cosine', re.embedding, @e) as semantic_similarity,
    @p.STDistance(geo_location) as geo_distance
from
    dbo.reviews r
inner join
    dbo.reviews_embeddings re on r.id = re.review_id
inner join |
    dbo.business b on r.business_id = b.id
where
    b.city = 'Seattle'
and
    @p.STDistance(b.geo_location) < 5000 -- 5 km
and
    r.stars ≥ 4
and
    b.reviews ≥ 30
and
    json_value(b.custom_attributes, '$.local_recommended') = 'true'
and
    vector_distance('cosine', re.embedding, @e) < 0.2
order by
    semantic_similarity desc

```



# RAG Pattern Explained

## Retrieval Augmented Generation



<https://learn.microsoft.com/en-us/azure/cosmos-db/gen-ai/rag-chatbot> (Cosmos DB)

<https://github.com/Azure-Samples/rag-postgres-openai-python> (PostgreSQL)

<https://github.com/Azure-Samples/azure-sql-db-chatbot> (Azure SQL)

# RAG Limitations

- RAG is fantastic when it comes to find *similar* content given a somehow approximate request.
- *Is it effective for all scenarios?*
- *What about very precise and direct request?*

For example: "Show me all the products in category VideoGames updated in the last 3 days"

Natural Language to SQL is *needed*

Navigate relationships can lead to better relevance in top results

# RAG Limitations

"Show me all the products in category VideoGames updated in the last 3 days"

```
SELECT
    *
FROM
    dbo.products p
INNER JOIN
    dbo.categories c ON p.category_id = c.id
WHERE
    c.Name = 'videogame'
AND
    p.last_updated >= DATEADD(day, -3, SYSDATETIME())
```

# AI Agents and Agentic RAG

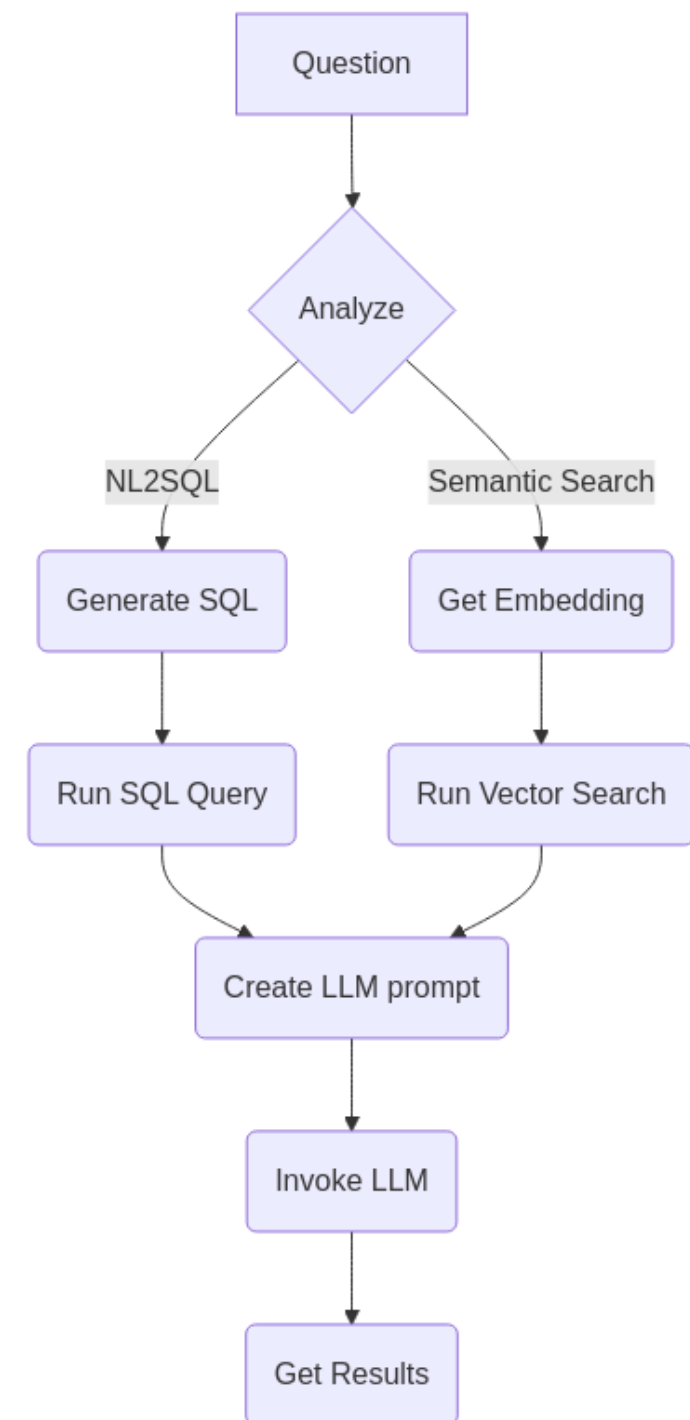
Now things are a bit more complex, and the process is more suitable for *agents*

Orchestrator agent: understand if the questions needs NL2SQL or not

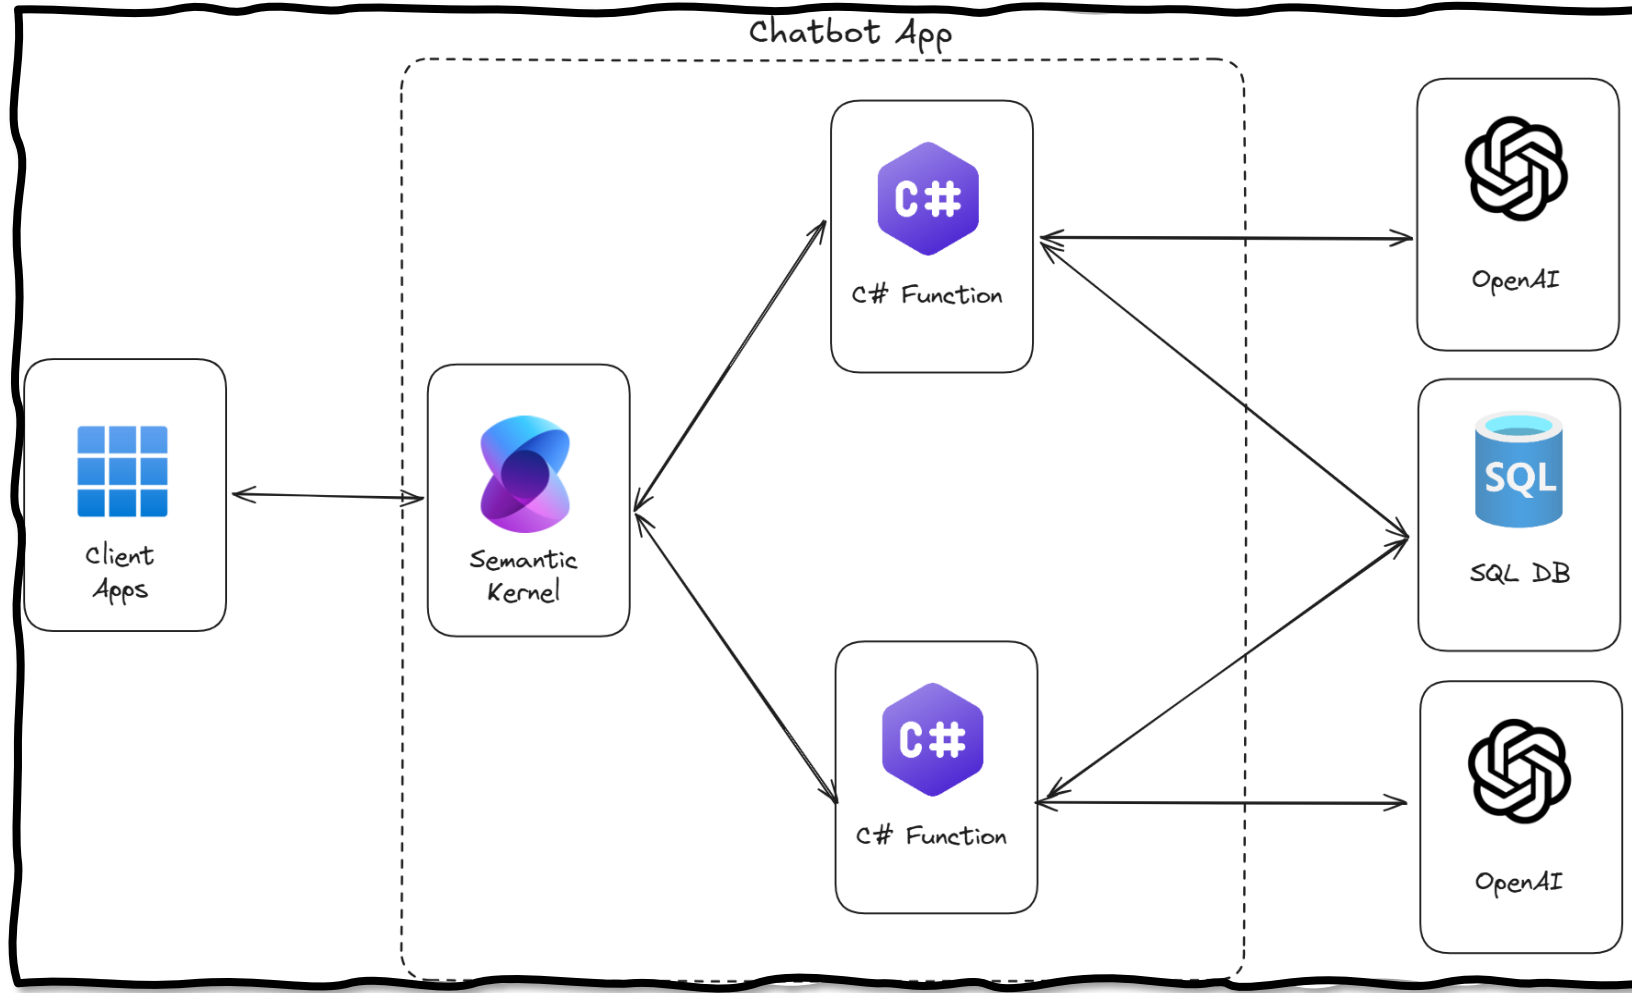
If yes, a specialized database agent (who has been informed of the database schema) will  
generate and execute the query  
Or just run the semantic search as usual

Use the output of the previous agent to build the LLM prompt

Invoke LLM and get results



# End-to-End Application (Agents with C#)




<https://github.com/Azure-Samples/azure-sql-db-chat-sk>





# CosmosAIGraph Architecture

## Development Environment

 python™  
fastapi  
uvicorn  
rdflib  
pymongo  
openai  
semantic-kernel  
  
Docker  
docker compose

## Container Registry



docker images (2)

graph



web



**Cosmos DB vCore**  
Libraries (domain)  
sessions  
prompts  
completions  
optional vector search



https

## Azure Container App (ACA)

Two Microservices – web and graph  
Scale and optimize each individually  
In-memory RDF graph. App in own Vnet.

az CLI  
program  
w/Bicep

deploy

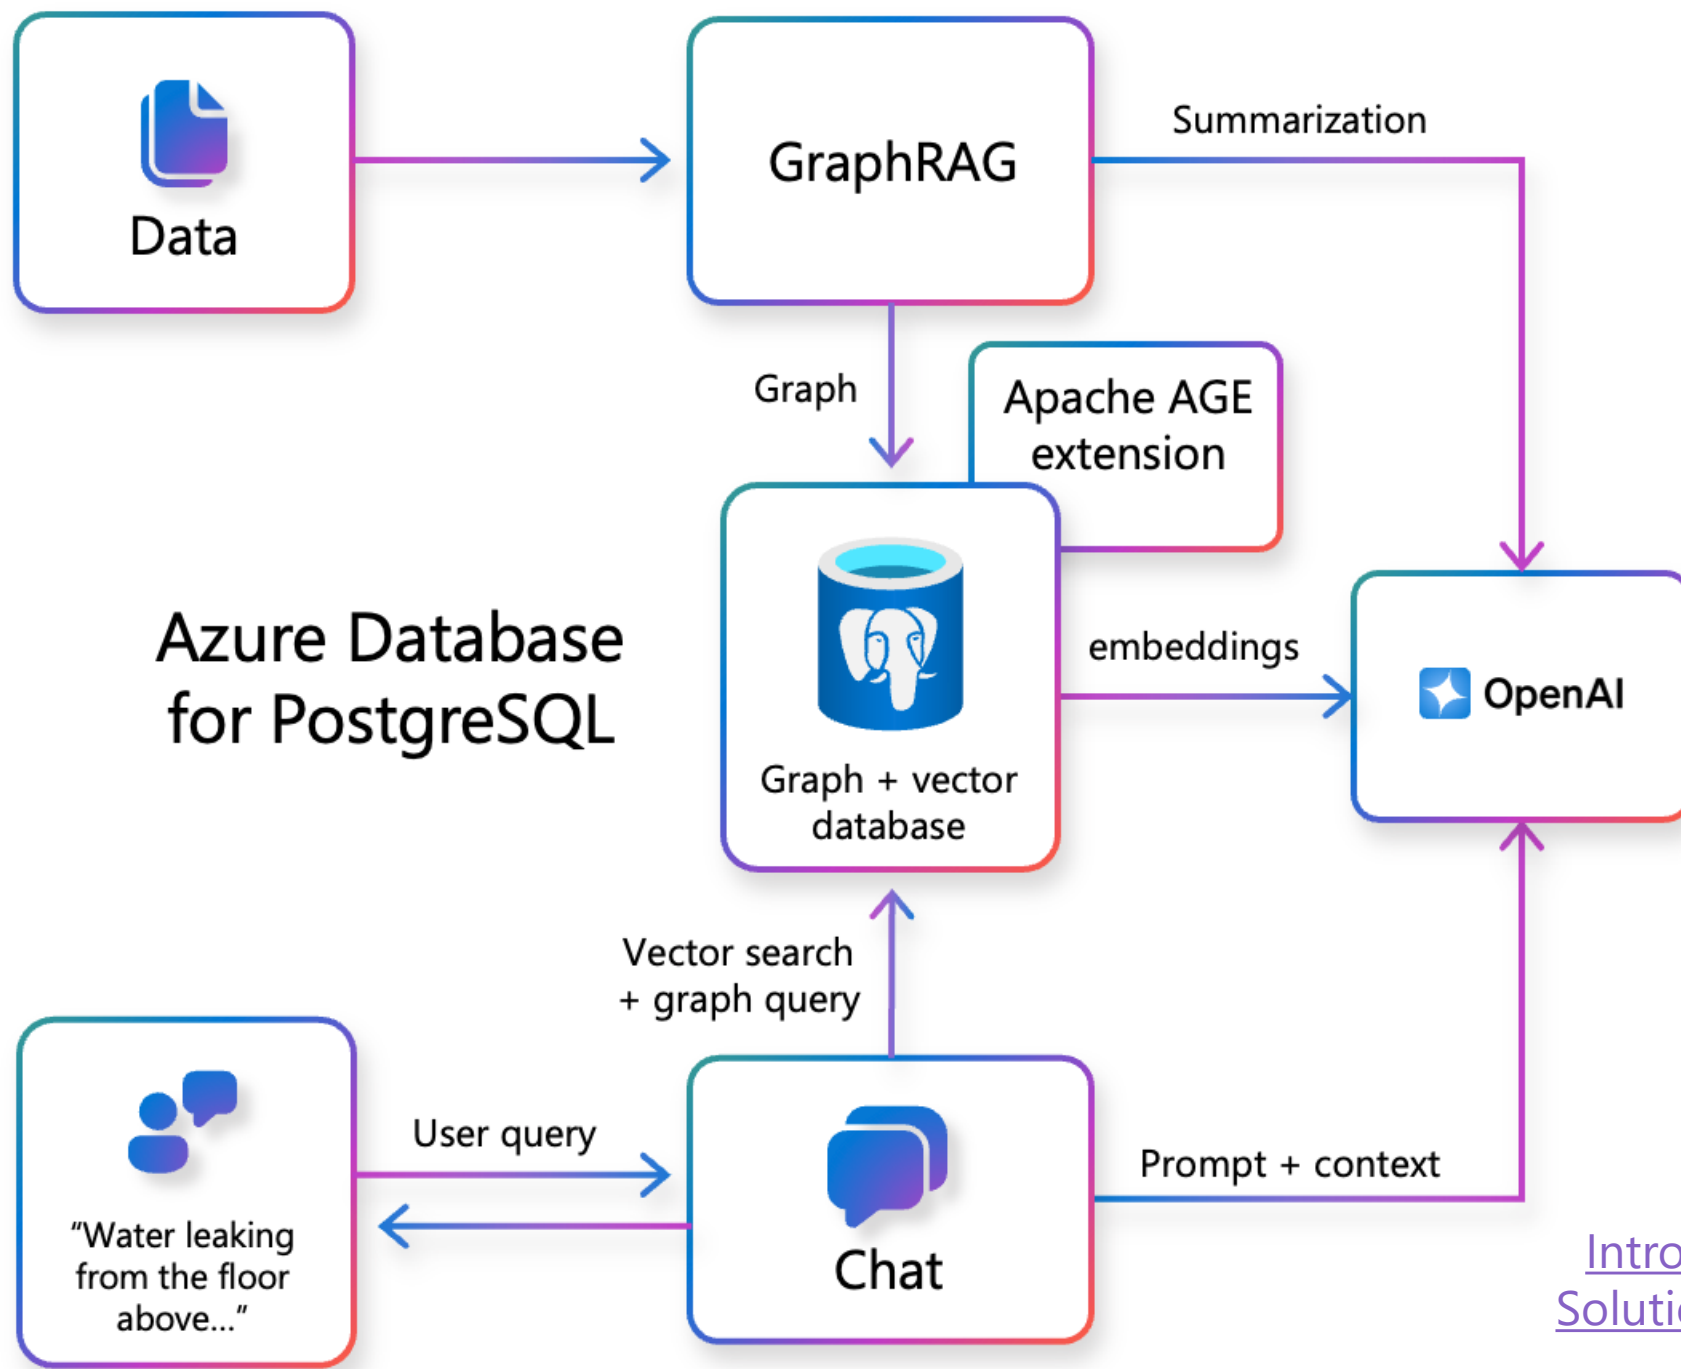


**Azure OpenAI**  
gpt-4  
completions  
text-embedding-ada-002  
embeddings

<https://github.com/cjoakim/CosmosAIGraph>

[AI knowledge graphs | Microsoft Learn](#)

<https://aka.ms/graphrag> (MSR)



Introducing the GraphRAG  
Solution for Azure Database  
for PostgreSQL

# AI and Agentic Agents



# Modern Agentic AI Insurance Sample



Orchestrator  
Agent

"I'm meeting with John Doe,  
What can you tell me about him?"




CRM Agent



Claims Agent



Payments  
Agent



What tools I can use to find  
customer's related data?

Who can help me?

# Modern Agentic AI Insurance Sample



"I'm meeting with John Doe,  
What can you tell me about him?"



Orchestrator  
Agent



CRM Agent



Claims Agent



Payments  
Agent

I know how to query CRM  
data

What tools I can use to find  
customer's related data?

Who can help me?

# Specialized Agent has an “expert-level” view of the data

The database schema is the following:


```
// this table contains customer information
CREATE TABLE [dbo].[customers]
(
    [id] INT DEFAULT (NEXT VALUE FOR [dbo].[global_id]) NOT NULL,
    [first_name] NVARCHAR(100) NOT NULL,
    [last_name] NVARCHAR(100) NOT NULL,
    [address] NVARCHAR(100) NOT NULL,
    [city] NVARCHAR(100) NOT NULL,
    [state] NVARCHAR(100) NOT NULL,
    [zip] NVARCHAR(100) NOT NULL,
    [country] NVARCHAR(100) NOT NULL,
    [details] JSON NULL, -- make sure to cast to NVARCHAR(MAX) before using it in a query w
    PRIMARY KEY NONCLUSTERED ([id] ASC)
)
```

the [details] column contains JSON data with the following structure:

active-policies: [string...string] other type of policies the customer has (life, health, c



# Modern Agentic AI Insurance Sample



"I'm meeting with John Doe,  
What can you tell me about him?"



Orchestrator  
Agent


I know how to work with  
claims data




CRM Agent



Claims Agent



Payments  
Agent



What tools I can use to find  
customer's related data?

Who can help me?

# Modern Agentic AI Insurance Sample



Orchestrator Agent

"I'm meeting with John Doe,  
What can you tell me about him?"



CRM Agent



Claims Agent



Payments Agent

I can work with payment data

What tools I can use to find  
customer's related data?


Who can help me?

# Modern Agentic AI Insurance Sample

Get details about "John Doe"



Orchestrator Agent



CRM Agent



Claims Agent



Payments Agent


```
SELECT * FROM pass.customers  
WHERE first_name = 'John' AND  
last_name = 'Doe'
```



"I'm meeting with John Doe,  
What can you tell me about him?"

# Modern Agentic AI Insurance Sample

John Doe: ID = 123, Address = XYZ, etc..



Orchestrator Agent

"I'm meeting with John Doe,  
What can you tell me about him?"



CRM Agent



ID = 123, Address = XYZ, etc..



Claims Agent



Payments Agent

# Modern Agentic AI Insurance Sample

John Doe: ID = 123, Address = XYZ, etc.



"I'm meeting with John Doe,  
What can you tell me about him?"



Orchestrator  
Agent



Claims Agent



```
SELECT * FROM dbo.claims  
WHERE CustomerId = 123
```



Payments  
Agent



CRM Agent

# Modern Agentic AI Insurance Sample

John Doe: ID = 123, Address = XYZ, etc.  
Claims: XYZ234, XYZ999



Orchestrator Agent

"I'm meeting with John Doe,  
What can you tell me about him?"

John Doe: ID = 123, Address = XYZ, etc.  
Claims: XYZ234, XYZ999  
Payments: (XYZ234, \$100, 2025-02-01)



**"I authorize payment for claim XYZ999"**



CRM Agent



Claims Agent



Payments Agent



```
SELECT * FROM dbo.payments  
WHERE CustomerId = 123 and  
ClaimId in ('XYZ234', 'XYZ999')
```



# What are agents?

AI designed to perform a task

Tasks can vary in level of complexity and capabilities depending on your need

Simple



## Generation

**Generate** summaries, images, audio, and more with an AI model and inputs.



## Retrieval

**Retrieve information** from grounding data, reason, summarize, and answer user questions



## Action

**Take actions** to automate workflows, and replace repetitive tasks for users

Advanced


# Modern Agentic AI Insurance Sample



"I'm meeting with John Doe,  
What can you tell me about him?"



Orchestrator  
Agent



CRM Agent



Claims Agent

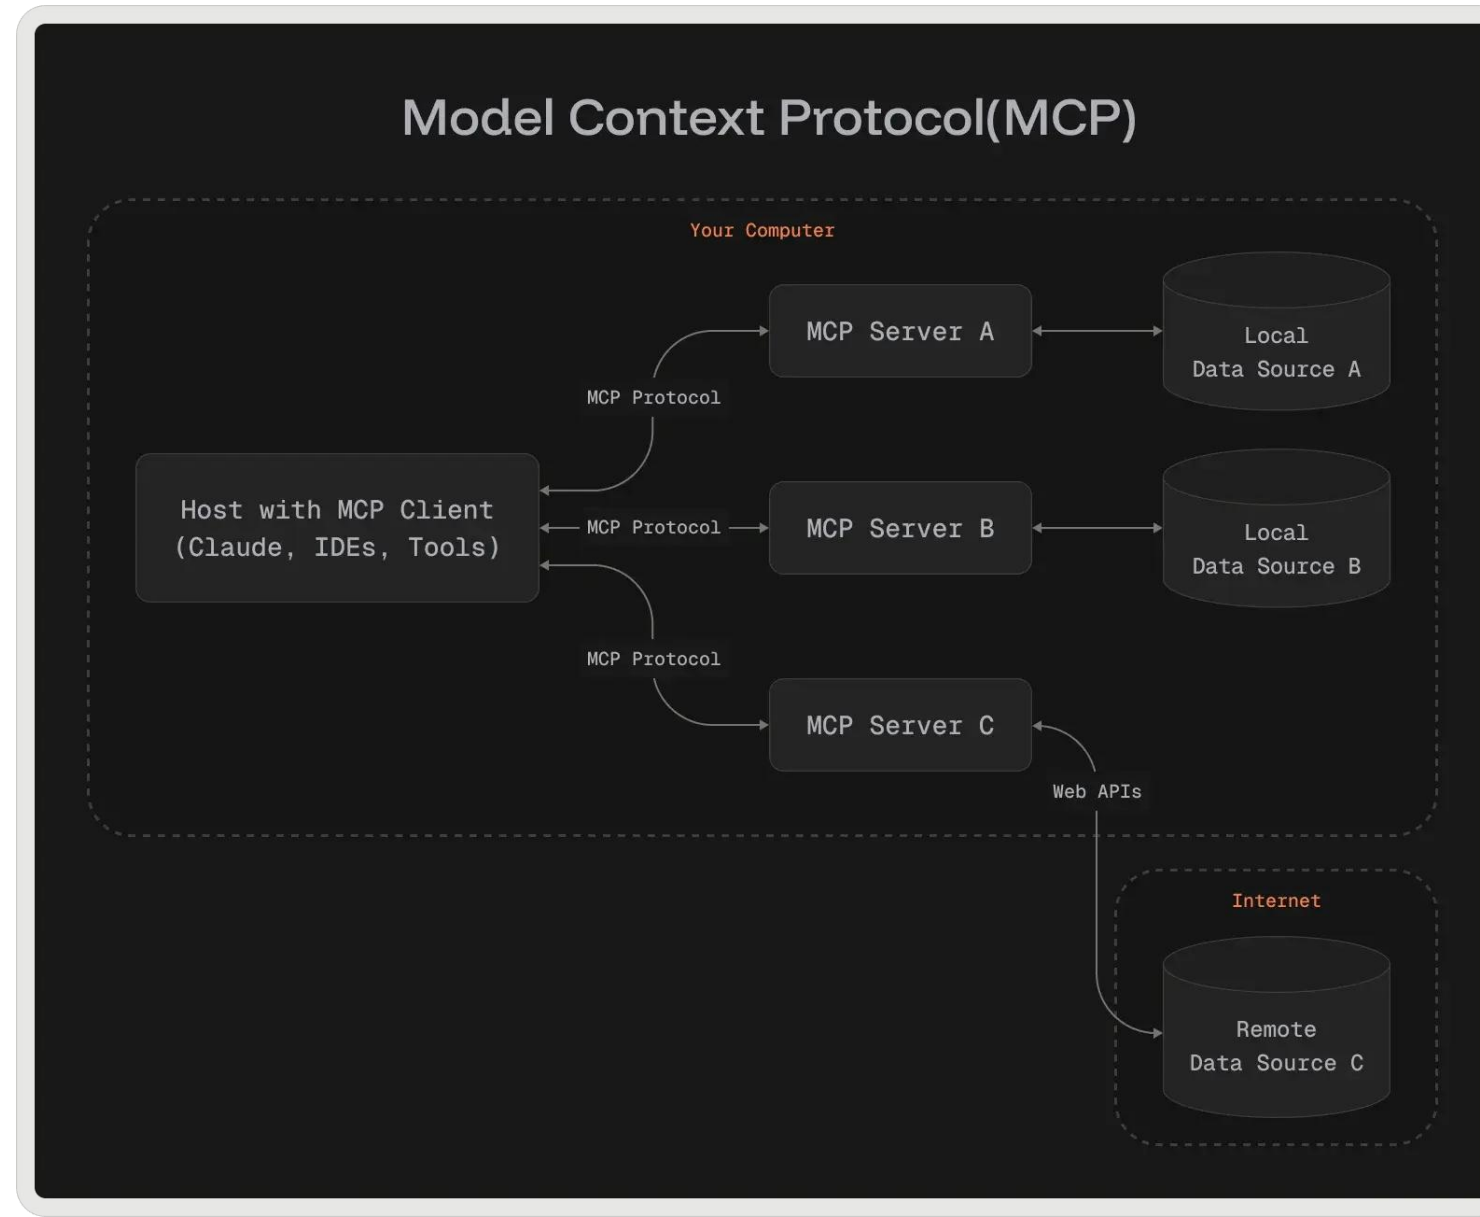


Payments  
Agent



# MCP – Model Context Protocol

The Model Context Protocol (MCP) is an **open standard** designed to **bridge AI applications** with external systems and data sources, ensuring secure interactions. It simplifies **how models interact** with various tools and databases through a single, secure interface, enabling them to **invoke functions and fetch data** efficiently



# Official Microsoft MCP Tools

- Official GitHub MCP Server [[Demo](#)]
- Azure SDK MCP Server [[Demo](#)]
- Azure AI Agent Service [[Demo](#)] [[GitHub](#)] [[Blog](#)]
- Playwright MCP Server [[GitHub](#)]
- Azure CosmosDB MCP Server [[GitHub](#)]
- Foundry Windows Actions Tool [[Demo](#)]

# References – Azure Cosmos DB

- <https://github.com/Azure-Samples/chat-with-your-data-solution-accelerator>
- <https://azurecosmosdb.github.io/gallery/>
- <https://github.com/Azure/document-vector-pipeline>
- <https://github.com/Azure-Samples/rag-postgres-openai-python/>
- <https://github.com/Azure-Samples/azure-postgres-pgvector-python>
- <https://devblogs.microsoft.com/cosmosdb/azure-cosmos-db-vector-search-with-diskann-part-1-full-space-search/>
- <https://devblogs.microsoft.com/cosmosdb/sharded-diskann-focused-vector-search-for-better-performance-and-lower-cost/>
- <https://devblogs.microsoft.com/cosmosdb/new-vector-search-full-text-search-and-hybrid-search-features-in-azure-cosmos-db-for-nosql/>

# References – Azure SQL

- [Demos from Fabric Conference workshop](#)
- [Semantic Kernel / Agentic RAG demo](#)
- [Vectors demo repo](#)
- [Chat with your data](#)
- [SQL Database Vector Search Sample](#)
- [Migrate and Modernize](#)



# References – Azure PostgreSQL

- [Introducing the GraphRAG Solution for Azure Database for PostgreSQL](#)
- <https://github.com/Azure-Samples/rag-postgres-openai-python>
- [Introducing the Semantic Ranking Solution for Azure Database for PostgreSQL](#)
- [Build AI Apps with Azure Database for PostgreSQL](#)
- <https://github.com/mcaps-microsoft/csa-cto>

# Thank you!

- [dandy.weyn@microsoft.com](mailto:dandy.weyn@microsoft.com)
- [luciano.moreira@microsoft.com](mailto:luciano.moreira@microsoft.com)
- [mehasodhi@microsoft.com](mailto:mehasodhi@microsoft.com)
- [prmadi@microsoft.com](mailto:prmadi@microsoft.com)